

6.033 Deep Dive 2.1

UNIX FILE SYSTEM LAYERING AND NAMING

Americo De Filippo

May 16, 2023

1 The UNIX file system

Unix is an operating system that is at the foundation of many others today, was first developed at Bell Telephone for the Digital Equipment Corporation PDP. Today instead is at the base of systems like GNU/Linux, which takes a lot from the standard UNIX system. Others systems like macOS are based on UNIX.

1.1 Application programming Interface for the UNIX File System

The UNIX file system provides an API for handling the files, the user can create, choose named files, create directory, create links and so on...

But the system has they have some way to organize all the files, especially those who are not created by the user. In the following sections we are going to have an understanding of the UNIX file system, moving our way up into all the layers of the file system. A good graphical representation is shown in the following figure.

Layer	Purpose	
Symbolic link layer	Integrate multiple file systems with symbolic links.	↑ user-oriented names
Absolute path name layer	Provide a root for the naming hierarchies.	
Path name layer	Organize files into naming hierarchies.	
File name layer	Provide human-oriented names for files.	machine-user interface
Inode number layer	Provide machine-oriented names for files.	↑ machine-oriented names
File layer	Organize blocks into files.	
Block layer	Identify disk blocks.	

Figure 1: The Naming layers of the UNIX file system

1.2 The Block Layer

The block layer is the lower one and for this reason the closer to the machine, in fact the Block layer *is* the layer which comprehend the physical stuff. In this layer we have the implementation of the physical stuff like hard drives that have to store all the *durable* information about the system. Usually the main hard drive is organized as shown in figure 2.



Figure 2: A simple disk layout

The first thing that we noticed is the naming, we have a number naming for each the memory block (also called address). The first element usually is where are store the information for booting the operating system, in the super block instead contains the size of the bitmap (for having an id of the occupied and the free file blocks) and the inode table in blocks (as explained next). Another important feature of the hard disk is the block size. The size of the section of the hard disk represent a problem for the operating system, because has to be made some trade-offs for having a good system. The usual problem could be that a small size can represent small files without wasting much memory, but could have problem with the handling of large size files cause has to find a lot of blocks that are difficult to organize into the memory (sequential blocks). For big size blocks instead the issue is with the wasting of memory that could be taken place when allocation small size file in a large size block.

1.3 The File Layer

The file layer is being implemented for handle large data that doesn't fit inside a single block of the memory. A file is a ordered flow of bytes, this flow of data has to be stored in blocks of the memory, that are much more difficult to be handled for a system, for the simple reason that it has to remember which blocks of the disk belongs to which file. For this reason the UNIX file system use the *inode*. An inode is a simple structure that stores the block numbers that constitute the file and the size in bytes of the it. With this information we can figure out a lot of information about the memory disposition of a single file.

1.4 The Inode Number Layer

Instead of passing table around the system we can have a compact representation of them, that will enable us to have a much simpler access. We will store all the inodes into a table (a regular table indexed that will allow a really simple access) then we are going to store her at the beginning of the disk so that with these three layers of naming we can with a simple procedure access at the index of the block of a file, or have just an high level overview of the system.

1.5 The File Name Layer

For people numbers are not so friendly to be used for the purpose of naming something instead they are more likely to use names = strings, for making naming. But the computer has (for memory scopes) name files with numbers (addresses). By default, the UNIX file system provides the LOOKUP(filename, dir) procedure, that searches for the file (string) inside a directory. But what is a directory? A directory is a collection of files, and for the UNIX file system the *inode* specify if it's a regular file or a directory. Usually a directory stores the map from the filename to inode numbers like that:

File name	Inode number
program	10
paper	12

Figure 3: The graphical representation of a directory in UNIX

1.6 The Path Name Layer

Another useful method for organize data is the through the naming of the directory. In UNIX a path structure is represented like this: project/file, where project is a directory with inside a file. The previous structure is called a tree of hierarchy that enables us to manage our system more fully. When we want from a path access let's say to the block number of a file we just make a recursive function that goes to branch of our tree till it reaches the file or the directory (in this case we will return a map from filename to blocks numbers).

1.7 Links

The Links in a system are of fundamental importance, they allow the system to become much more customizable and organized. The feature that implements the link is the procedure LINK(path, file), this procedure allows you to create a link between the path and a file so to squash everything. The UNIX system will deny the linking of a path to a directory for the possible creation of cycles like the following: LINK(a/b/c, a) this will create a cycle from c to a. Moreover, every inode stores a counter a reference to itself. In this way the deleting function could be taken place only when the counter is at zero (the counter decrease with the UNLINK(path) procedure).

1.7.1 Soft Links

In Linux we can see a simple implemetation of soft link, we usually refer to them as simple shortcuts to a file, in fact they preserv their own inode, because they

are a file which points to another file in the memory but with their own inode

1.7.2 Hard Links

Hard links on the other hand are a bit more different, when we define an Hard Link in UNIX we are basically creating as second pointer to a file, but despite the soft side he owns the file. This concept of owning means that the file will have had an increase in reference pointer, meaning that if a make an Hard link to let's say text.txt, then i delete text.txt, it actually still exist in the file which i have made the hard link to.

1.8 Renaming

We can implement the renaming practive via link and unlink file (here we mean Hard Links) the Renaming could be implemented in the following way:

LINK(from-name, name);

UNLINK(from-name);

In this practive we are creating an Hard link to the file, increasing the reference number of 1 then we are decreasing the reference number removing the first hard link (which we have created when creating the file).

1.9 The Absolute Path Name Layer

At this moment we are able to create some kind of user-defined heierarchy but we are able to share data or programs between users. So we need some kind of space where all the data are in a context which is common to all users. This space is implemented before the user-name directory, this space starts with the root directory. This way of rappresenting the path is called absolute path, they are usueally of the form: /usr/bin/..., the first folder (which in this case in usr) has a link to itself and if you try to perform a '..' operation you will discover that she is her own parents directory. A good way of viewing this is by look at this rappresentation of how we are going to locate the content of a file in the memory just by knowing the path. For example, to find the blocks corresponding to the file “/programs/pong.c” with the information in Figure 4, we start by finding the inode table, which starts at a block number (block 4 in our example) stored in the super block . From there we locate the root inode (which is known to be inode number 1). The root inode con- tains the block numbers that in turn contain the blocks of the root directory; in the figure the root starts in block number 14. Block 14 lists the entries in the root direc- tory: “programs” is named by inode number 7. The inode table says that data for inode number 7 starts in block number 23, which contains the contents of the “ programs” directory. The file “pong.c” is named by inode number 9. Referring once more to the inode table, to see where inode 9 is stored, we see that the data corresponding to inode 9 starts in block number 61. In short, directories and files are carefully laid out so that all information can be found by starting from the well-known location of the root inode.

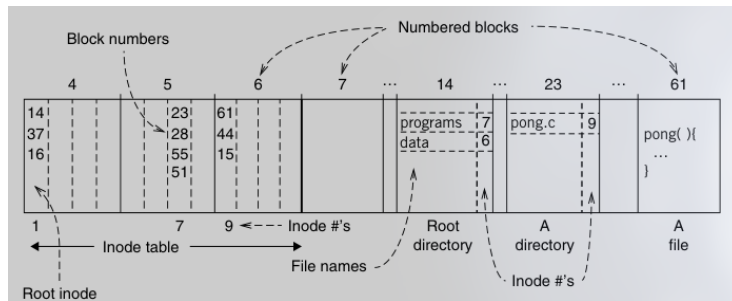


Figure 4: Memory representation of a file and hierarchy

1.10 The Symbolic Link Layer

The UNIX file system also enables the user to create and manage file system on others disks, this practice is supported by the mount and unmount commands. Let's say we are mounting the `/dev/df1` onto `/flash`, this means that all the file system of the first disk is now mounted into the flash directory. Usually the mount does not remain after shutting down the system.

1.11 Implementing the File System API

We have other important things to implement for making our system complete. We have seen that almost all the devices are implemented via files, the basic procedure that the API has to provide are `OPEN`, `READ`, `WRITE`, `CLOSE`. This procedure has a lot of important jobs; they have to provide a lot of information when using a file, like the last time editing, the cursor block position etc... Another important aspect of the file system that has to do with files, is when a thread inherits a file, and it will have all the current information of the process that has generated it. Another interesting property is the dynamic storage of the files. Usually the system stores the last used or most frequently used files in the cache (for memory performance) memory which is the one that it first searches when finding a file.